

```
import tensorflow as tf
import numpy as np

class classifier:
    def __init__(self,
                  learning_rate,
                  hidden_size,
                  batch_size,
                  max_time,
                  embeddings,
                  global_step):

        self.batch_size = batch_size
        self.learning_rate = tf.train.exponential_decay(
            learning_rate, global_step, 500, 0.8, staircase=True)
        self.hidden_size = hidden_size
        self.batch_size = batch_size
        self.max_time = max_time
        self.embeddings = tf.constant(embeddings, name = "emb")
        self.optimizer = tf.train.AdamOptimizer(self.learning_rate)
        self.global_step = global_step

        tf.summary.scalar('Learning Rate', self.learning_rate)

        self.make()

        self.saver = tf.train.Saver(max_to_keep = 200)

    def make(self):
        self.sequence_length = tf.placeholder(tf.float32, shape = (self.
batch_size))
        self.inputs = tf.placeholder(tf.int32, shape = (self.batch_size,
self.max_time))
        self.embed = tf.nn.embedding_lookup(self.embeddings, self.inputs
)
        self.targets = tf.placeholder(tf.int64, shape = (self.batch_size
,2))
        self.keep_prob = tf.placeholder(tf.float32, shape = None)

        cell = tf.nn.rnn_cell.LSTMCell(
            num_units = self.hidden_size,
            use_peepholes=True)

        dropout_cell = tf.contrib.rnn.DropoutWrapper(
            cell = cell,
            output_keep_prob = self.keep_prob
        )

        self.output, state = tf.nn.dynamic_rnn(
            cell = dropout_cell,
            inputs = self.embed,
            sequence_length = self.sequence_length,
            initial_state = cell.zero_state(self.batch_size, tf.float32))

        self.sum = tf.reduce_sum(self.output, axis = 1)
        self.mean = tf.divide(self.sum, tf.expand_dims(self.sequence_len
gth,1))
```

```
self.logits = tf.contrib.layers.fully_connected(self.mean, 2, activation_fn = None)
#self.logits = self.mean
self.probability = tf.nn.softmax(self.logits)
self.decision = tf.argmax(self.probability, axis=1)
self.actual = tf.argmax(self.targets, axis=1)
self.probe = self.decision
self.metric_accuracy, self.update_accuracy = \
    tf.metrics.accuracy(self.actual, self.decision)

self.xent = tf.nn.softmax_cross_entropy_with_logits(
    labels = self.targets,
    logits = self.logits)

self.loss = tf.reduce_mean(self.xent)
# warning: changed logits to probability - may be numerically unstable
self.pos_grad = tf.gradients(self.probability[0][0], self.embed)
self.neg_grad = tf.gradients(self.probability[0][1], self.embed)
self.logit_grad = tf.gradients(self.logits[0][0], self.embed)
self.updates = self.optimizer.minimize(self.loss, global_step = self.global_step)

tf.summary.scalar('Metric Accuracy', self.update_accuracy)
tf.summary.scalar('Loss', self.loss)

self.merged = tf.summary.merge_all()

def infer_dpg(self, sess, rv):
    decision, probability, grad = \
        sess.run([self.decision, self.probability, self.pos_grad],
            feed_dict = \
                {self.inputs: rv.index_vector,
                 self.targets: rv.targets,
                 self.sequence_length: [rv.length],
                 self.keep_prob: 1.0})

    return decision, probability, grad

def infer_rep_dpg(self, sess, rv, word_index):
    index_matrix = np.tile(rv.index_vector, (rv.length, 1))
    np.fill_diagonal(index_matrix, word_index)
    target_matrix = np.tile(rv.targets, (rv.length, 1))

    decision, probability = \
        sess.run([self.decision, self.probability],
            feed_dict = \
                {self.inputs: index_matrix,
                 self.targets: target_matrix,
                 self.sequence_length: [rv.length]*rv.length,
                 self.keep_prob: 1.0})

    return decision, probability

def infer_batched_prob(self, sess, rv, word_index, per_batch, top_idx):
    num_top = self.batch_size//per_batch
```

```

index_matrix = np.tile(rv.index_vector, (self.batch_size,1))
target_matrix = np.tile(rv.targets, (self.batch_size,1))
n = 0
for idx in top_idx:
    for i in range(per_batch):
        index_matrix[n,idx] = word_index + i
        n = n + 1

decision, probability, grad = \
    sess.run([self.decision,self.probability,self.pos_grad],
        feed_dict = \
            {self.inputs:index_matrix,
             self.targets:target_matrix,
             self.sequence_length:[rv.length]*self.batch_size,
             self.keep_prob:1.0})

return decision, probability, grad

# inserts all words at index 'insert_location'
def infer_insert(self, sess, rv, insert_location, divs, top_k):
    per = rv.length // divs
    end = per*divs
    new = per+1
    index_matrix = np.zeros((self.batch_size,new),'int')
    wim = np.reshape(rv.index_vector[0,0:end], (per,divs),'F')
    wim = np.insert(wim,insert_location,0,axis=0)
    wim = np.reshape(wim, (1,end+divs),'F')
    target_matrix = np.tile(rv.targets, (self.batch_size,1))
    for i in range(divs):
        start = i*top_k; end = (i+1)*top_k
        index_matrix[start:end,:] = np.tile(wim[0,new*i:new*(i+1)]
, (top_k,1))
    #index_matrix = np.tile(wim, (self.batch_size//divs,1))
    index_matrix[:,insert_location] = \
        np.arange(self.batch_size) % (self.batch_size//divs)
    d,p,g = sess.run([self.decision, self.probability, self.pos_g
rad],
        feed_dict = \
            {
                self.inputs: index_matrix,
                self.targets: target_matrix,
                self.sequence_length: [new] * self.batch_size,
                self.keep_prob: 1.0
            })
    return d,p,g,index_matrix

def infer_swap(self, sess, rv, swap_location, divs, top_k):
    per = rv.length // divs
    end = per*divs
    index_matrix = np.zeros((self.batch_size,per),'int')
    wim = np.copy(np.reshape(rv.index_vector[0,0:end], (per,divs),'
F'))
    wim[swap_location,:] = 0
    wim = np.reshape(wim, (1,end),'F')
    target_matrix = np.tile(rv.targets, (self.batch_size,1))
    replacement_vector = np.random.choice(np.arange(10000),size=se
lf.batch_size)

```

```

    for i in range(divs):
        start = i*top_k; end = (i+1)*top_k
        index_matrix[start:end,:] = np.tile(wim[0,per*i:per*(i+1)]
, (top_k,1))
        index_matrix[:,swap_location] = replacement_vector
        #np.arange(self.batch_size) % (self.batch_size//divs)
        d,p,g = sess.run( [self.decision, self.probability, self.pos_g
rad],
            feed_dict = \
                {
                    self.inputs: index_matrix,
                    self.targets: target_matrix,
                    self.sequence_length: [per] * self.batch_size,
                    self.keep_prob: 1.0
                })
    return d,p,g,index_matrix,replacement_vector

def infer_window(self, sess, rv, word_index, window_size):
    w = window_size
    L = rv.length
    n = word_index
    n1 = n - w//2; n2 = n1 + w
    if n1 >= 0 and n2 <= L:
        i1 = n1; i2 = n2
    elif n1 < 0 and n2 <= L:
        i1 = 0; i2 = w
    elif n1 >=0 and n2 > L:
        i1 = L-w; i2 = L
    else:
        i1 = 0; i2 = rv.length
    index_matrix = np.tile(rv.index_vector[0,i1:i2], (10000,1))
    index_matrix[:,word_index-i1] = np.arange(10000)
    target_matrix = np.tile(rv.targets, (self.batch_size,1))

    d,p,g = sess.run( [self.decision, self.probability, self.pos_g
rad],
        feed_dict = \
            {
                self.inputs: index_matrix,
                self.targets: target_matrix,
                self.sequence_length: [w] * self.batch_size,
                self.keep_prob: 1.0
            })
    return d,p,g

def infer_multi(self, sess, rv, ii, K):
    # K is the list of source word indices to be replaced
    # batch size is fixed at 10,000 with number of
    # destination words = 10,000 / len(K)
    N = self.batch_size // K.size
    index_matrix = np.tile(rv.index_vector[0,:], (self.batch_size,1
))
    c = 0
    for k in list(K):
        index_matrix[c*N:(c+1)*N,k] = ii[np.arange(N),k]
        c = c + 1
    target_matrix = np.tile(rv.targets, (self.batch_size,1))

```

```
d,p,g = sess.run( [self.decision, self.probability, self.pos_g
rad],
    feed_dict = \
        {
            self.inputs: index_matrix,
            self.targets: target_matrix,
            self.sequence_length: [rv.length] * self.batch_size,
            self.keep_prob: 1.0
        })
return d,p,g
```